

Assignment 1 - Part B2: Working with you own data

In this assignment you will create you own dataset for classification. You will explore which ML algorithms are best to classify this and you will present your best solution.

Generating your dataset:

For this assignment you will create your own dataset of motions that you collect with an Accelerometer and Gyroscope. For this you can use your phone as a sensor. To be able to collect your data you can best use an app called [phyphox](#), this is a free app available in app stores. This app can be configured to acces your sensordata, sample it as given frequency's. you can set it up te have experiment timeslots, and the data with a timestamp can be exported to a needed output format.



When you installed the app you can setup a custum experiment by clicking on the + button. Define an experiment name, sample frequency and activate the Accelerometer and Gyroscope. Your custom experiment will be added, you can run it pressing the play button and you will see sensor motion. Pressing the tree dots (...) lets you define timed runs, remote access and exporting data.

steps

With your own generated dataset the similar sequence of steps should be taken to train your model.

These are the generic steps to be taken

1. Frame the problem and look at the big picture.
2. Get the data.
3. Explore the data to gain insights.
4. Prepare the data to better expose the underlying data patterns to Machine Learning algorithms
5. Explore many different models and short-list the best ones.
6. Fine-tune your models and combine them into a great solution.
7. Present your solution.
8. Launch, monitor, and maintain your system.
9. Additional Questions

In the Notebook this structure is used for dividing the different steps, so make sure you do the implementation and analysis at these location in the notebook.

You may add additional code blocks, but keep the separation of the given structure.

At the end of each block summarize / comment / conclude your current step in the given textblocks.

Hints

Additional info can be found in the tips and trick document

Felix Douven, Kasper Walraven, Yosha op het Veld

1. Frame the problem and look at the big picture

Describe the problem at hand and explain your approach

The goal is to create a model which is based on the "Train" data that can correctly predict which activity a subject is undertaking.

The data will be analysed, and checked for any errors such as missing data or duplicate data. Multiple algorithms will be tested and plotted to determine the most successful algorithm(s). This model consists of the most successful algorithm(s) tested.

2. Get the data.

Initialize the system, get all needed libraries, retrieve the data and import it

NOTE: You can download the dataset directly from github, see Tips & Tricks

```
In [1]: from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
import os
import time
from sklearn import model_selection
from sklearn.model_selection import cross_validate
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from pandas.plotting import scatter_matrix

from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import pandas as pd
from sklearn.neural_network import MLPClassifier
```

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.gaussian_process.kernels import RBF
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from pandas.plotting import scatter_matrix
import matplotlib.pyplot as plt
from sklearn.datasets import load_wine
from warnings import simplefilter
from sklearn.model_selection import cross_val_score
import numpy as np
from sklearn.metrics import accuracy_score
from sklearn.metrics import matthews_corrcoef
from sklearn.metrics import f1_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import LogisticRegression
import warnings
import matplotlib.pyplot as plt
import seaborn as sns
from plotly.offline import init_notebook_mode, iplot
import plotly.graph_objs as go
from matplotlib.colors import rgb2hex
from matplotlib.cm import get_cmap
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.figure_factory as ff
from mpl_toolkits import mplot3d
from pylab import rcParams
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression, RidgeClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, f1_
from sklearn.model_selection import RandomizedSearchCV
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, StackingClassifier, GradientBoostin
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
import zipfile
import pandas as pd
import sklearn
import matplotlib.pyplot as plt
from sklearn import model_selection
from sklearn.model_selection import cross_validate
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from pandas.plotting import scatter_matrix
from sklearn.model_selection import GridSearchCV

```

Here we will import all the libraries that are needed in the program

```
In [2]: test = pd.read_csv("test.csv")
train= pd.read_csv("train.csv")
data = pd.concat([test, train])

print(data)
print(train)
```

	Time (s)	Acceleration x (m/s^2)	Acceleration y (m/s^2)	\
0	0.152	-2.59	4.11	
1	0.202	-1.87	4.41	
2	0.252	-2.03	3.78	
3	0.302	-1.83	3.82	
4	0.352	-1.17	2.31	
...	
21091	105.000	8.09	-1.47	
21092	105.000	8.08	-1.46	
21093	105.000	8.07	-1.45	
21094	NaN	NaN	NaN	
21095	NaN	NaN	NaN	

	Acceleration z (m/s^2)	Gyroscope x (rad/s)	Gyroscope y (rad/s)	\
0	8.79	-0.525000	-0.393000	
1	8.58	-0.645000	-0.235000	
2	8.92	-0.765000	-0.076200	
3	9.06	-0.320000	-0.044600	
4	8.81	-0.328000	-0.148000	
...	
21091	5.13	-0.001860	0.001330	
21092	5.14	0.000799	0.000133	
21093	5.15	0.002930	0.000732	
21094	NaN	NaN	NaN	
21095	NaN	NaN	NaN	

	Gyroscope z (rad/s)	subject	activity
0	0.08300	3.0	cycling
1	0.29200	3.0	cycling
2	0.38700	3.0	cycling
3	0.68200	3.0	cycling
4	0.50100	3.0	cycling
...
21091	-0.00413	2.0	sitting
21092	-0.00346	2.0	sitting
21093	-0.00226	2.0	sitting
21094	NaN	NaN	NaN
21095	NaN	NaN	NaN

[42381 rows x 9 columns]

	Time (s)	Acceleration x (m/s^2)	Acceleration y (m/s^2)	\
0	217.0	-4.75	-6.60	
1	217.0	-7.49	-7.82	
2	217.0	-8.96	-7.61	
3	217.0	-7.71	-7.38	
4	217.0	-10.10	-8.47	
...	
21091	105.0	8.09	-1.47	
21092	105.0	8.08	-1.46	
21093	105.0	8.07	-1.45	
21094	NaN	NaN	NaN	
21095	NaN	NaN	NaN	

	Acceleration z (m/s^2)	Gyroscope x (rad/s)	Gyroscope y (rad/s)	\
0	-1.500	0.560000	-0.400000	

1	1.820	0.590000	-0.726000
2	0.501	0.263000	-0.697000
3	3.060	-0.450000	0.231000
4	3.390	-0.721000	0.550000
...
21091	5.130	-0.001860	0.001330
21092	5.140	0.000799	0.000133
21093	5.150	0.002930	0.000732
21094	NaN	NaN	NaN
21095	NaN	NaN	NaN

	Gyroscope z (rad/s)	subject	activity
0	1.62000	1.0	cycling
1	2.08000	1.0	cycling
2	1.71000	1.0	cycling
3	0.91500	1.0	cycling
4	0.01400	1.0	cycling
...
21091	-0.00413	2.0	sitting
21092	-0.00346	2.0	sitting
21093	-0.00226	2.0	sitting
21094	NaN	NaN	NaN
21095	NaN	NaN	NaN

[21096 rows x 9 columns]

Importing CSV files from same directory and printing them

3. Explore the data to gain insights.

Explore the data in any possible way, visualize the results (if you have multiple plots of the same kind of data put them in one larger plot)

NOTE: You can visualize high-dimensional data in 2-d using T-distributed Stochastic Neighbor Embedding, see Tips & Tricks. (You can also visualize it in 3D, as described in the tutorial)

```
In [3]: print(f"Train Dataset Shape: {train.shape}")
print(f"Test Dataset Shape: {test.shape}")
print("-----")
print(f"Train Dataset Missing Data Counts: {train.isna().sum().sum()}")
print(f"Test Dataset Missing Data Counts: {test.isna().sum().sum()}")
print("-----")
print(f"Train Dataset Duplicate Data Counts: {train.duplicated().sum()}")
print(f"Test Dataset Duplicate Data Counts: {test.duplicated().sum()}")
```

```
Train Dataset Shape: (21096, 9)
Test Dataset Shape: (21285, 9)
```

```
-----
Train Dataset Missing Data Counts: 18
Test Dataset Missing Data Counts: 0
```

```
-----
Train Dataset Duplicate Data Counts: 1
Test Dataset Duplicate Data Counts: 0
```

Here you can see that we are missing 18 data. This are the last two rows of all the data. We will need to delete this before we can continue

```
In [4]: print("DataFrame after removing NaN values...", train.dropna())
```

```
train = train.dropna()
data = pd.concat([test, train])
```

```
DataFrame after removing NaN values...      Time (s)  Acceleration x (m/s^2)  Accelera
tion y (m/s^2)  \
0          217.0          -4.75          -6.60
1          217.0          -7.49          -7.82
2          217.0          -8.96          -7.61
3          217.0          -7.71          -7.38
4          217.0         -10.10          -8.47
...          ...          ...          ...
21089       105.0           8.08          -1.46
21090       105.0           8.09          -1.46
21091       105.0           8.09          -1.47
21092       105.0           8.08          -1.46
21093       105.0           8.07          -1.45

      Acceleration z (m/s^2)  Gyroscope x (rad/s)  Gyroscope y (rad/s)  \
0          -1.500          0.560000          -0.400000
1           1.820          0.590000          -0.726000
2           0.501          0.263000          -0.697000
3           3.060         -0.450000           0.231000
4           3.390         -0.721000           0.550000
...          ...          ...          ...
21089         5.140         -0.000200           0.001260
21090         5.120         -0.002460          -0.000333
21091         5.130         -0.001860           0.001330
21092         5.140          0.000799           0.000133
21093         5.150          0.002930           0.000732

      Gyroscope z (rad/s)  subject activity
0           1.62000         1.0  cycling
1           2.08000         1.0  cycling
2           1.71000         1.0  cycling
3           0.91500         1.0  cycling
4           0.01400         1.0  cycling
...          ...          ...          ...
21089         0.00153         2.0  sitting
21090        -0.00233         2.0  sitting
21091        -0.00413         2.0  sitting
21092        -0.00346         2.0  sitting
21093        -0.00226         2.0  sitting

[21094 rows x 9 columns]
```

In this printed data we dont see any missing data anymore.

```
In [5]: print(f"Train Dataset Shape: {train.shape}")
print(f"Test Dataset Shape: {test.shape}")
print("-----")
print(f"Train Dataset Missing Data Counts: {train.isna().sum().sum()}")
print(f"Test Dataset Missing Data Counts: {test.isna().sum().sum()}")
print("-----")
print(f"Train Dataset Duplicate Data Counts: {train.duplicated().sum()}")
print(f"Test Dataset Duplicate Data Counts: {test.duplicated().sum()}")
```

```
Train Dataset Shape: (21094, 9)
Test Dataset Shape: (21285, 9)
-----
Train Dataset Missing Data Counts: 0
Test Dataset Missing Data Counts: 0
-----
Train Dataset Duplicate Data Counts: 0
TestDataset Duplicate Data Counts: 0
```

This shows The shape of the data [columns, rows]

The "missing data counts" shows that no data is missing


The "duplicate data counts" shows that no data is the same

```
In [6]: plt.figure(figsize=(25, 5))
count_plot=sns.countplot(data=train, x='subject', hue='activity')
plt.gca().tick_params(axis='x')
plt.gca().tick_params(axis='y')
plt.xlabel( xlabel='Subjects')
plt.ylabel( ylabel='Activity Counts')
plt.legend(["cycling", "walking", "standing", "downstairs", "upstairs", "sitting"],bbox_
plt.title("Subjects Wise Activity Counts Train Set", fontsize=25, loc='left', pad=50)
plt.show()

plt.figure(figsize=(5, 5))
label_counts = train['activity'].value_counts()
colors = px.colors.qualitative.Plotly


graph = go.Bar(x=label_counts.index, y=label_counts.values, marker = dict(color = colors
layout = go.Layout(
    height=450, width=1100,
    title = 'Activity Counts Distribution Train Set',
    xaxis = dict(title = 'Activity', tickangle=0, showgrid=False),
    yaxis = dict(title = 'Count', showgrid=False),
    plot_bgcolor='#2d3035', paper_bgcolor='#2d3035',
    title_font=dict(size=25, color='#a5a7ab'),
    margin=dict(t=80, b=30, l=70, r=40),
    font=dict(color='#8a8d93'))
fig = go.Figure(data=[graph], layout = layout)
fig.update_traces(textfont=dict(color='#fff'), marker=dict(line=dict(color='#ffffff', wi
iplot(fig)
```

Subjects Wise Activity Counts Train Set



Activity Counts Distribution Train Set





<Figure size 500x500 with 0 Axes>

This plot shows how much data each activity contains. The amount of rows per activity equals to how much data is acquired for that activity.

```
In [7]: x_for_tsne = train.drop(['subject', 'activity'], axis=1)


tsne = TSNE(random_state = 42, n_components=2, verbose=1, perplexity=50, n_iter=1000).fit
plt.figure(figsize=(12,8))
sns.scatterplot(x =tsne[:, 0], y = tsne[:, 1], hue = train["activity"],palette="bright")
```

D:\programs\anaconda\lib\site-packages\sklearn\manifold_t_sne.py:780: FutureWarning:
The default initialization in TSNE will change from 'random' to 'pca' in 1.2.

D:\programs\anaconda\lib\site-packages\sklearn\manifold_t_sne.py:790: FutureWarning:
The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.

```
[t-SNE] Computing 151 nearest neighbors...
[t-SNE] Indexed 21094 samples in 0.043s...
[t-SNE] Computed neighbors for 21094 samples in 0.969s...
[t-SNE] Computed conditional probabilities for sample 1000 / 21094
[t-SNE] Computed conditional probabilities for sample 2000 / 21094
[t-SNE] Computed conditional probabilities for sample 3000 / 21094
[t-SNE] Computed conditional probabilities for sample 4000 / 21094
[t-SNE] Computed conditional probabilities for sample 5000 / 21094
[t-SNE] Computed conditional probabilities for sample 6000 / 21094
[t-SNE] Computed conditional probabilities for sample 7000 / 21094
[t-SNE] Computed conditional probabilities for sample 8000 / 21094
[t-SNE] Computed conditional probabilities for sample 9000 / 21094
[t-SNE] Computed conditional probabilities for sample 10000 / 21094
[t-SNE] Computed conditional probabilities for sample 11000 / 21094
[t-SNE] Computed conditional probabilities for sample 12000 / 21094
[t-SNE] Computed conditional probabilities for sample 13000 / 21094
[t-SNE] Computed conditional probabilities for sample 14000 / 21094
[t-SNE] Computed conditional probabilities for sample 15000 / 21094
[t-SNE] Computed conditional probabilities for sample 16000 / 21094
[t-SNE] Computed conditional probabilities for sample 17000 / 21094
[t-SNE] Computed conditional probabilities for sample 18000 / 21094
[t-SNE] Computed conditional probabilities for sample 19000 / 21094
[t-SNE] Computed conditional probabilities for sample 20000 / 21094
[t-SNE] Computed conditional probabilities for sample 21000 / 21094
[t-SNE] Computed conditional probabilities for sample 21094 / 21094
[t-SNE] Mean sigma: 1.280513
[t-SNE] KL divergence after 250 iterations with early exaggeration: 73.156456
[t-SNE] KL divergence after 1000 iterations: 1.001521
```

Out[7]: <AxesSubplot:>



In a scatterplot groups are easily identifiable except for upstairs and downstairs. there are no really clear areas for every activity but each activity does have some kind of line. So it should be possible to use the data.

```
In [8]: x_for_tsne = train.drop(['subject', 'activity'], axis=1)

tsne = TSNE(random_state = 42, n_components=3, verbose=1, perplexity=50, n_iter=1000).fit(x_for_tsne)

fig = px.scatter_3d(
    x = tsne[:, 0],
    y = tsne[:, 1],
    z = tsne[:, 2],
    color=train['activity']
)
fig.update_layout(
    title="Cluster Of Activities",
    title_font=dict(size=25, color='#a5a7ab'),
    font=dict(color='#8a8d93'),
    plot_bgcolor='#2d3035', paper_bgcolor='#2d3035',
    margin=dict(t=100, b=10, l=70, r=40),
)
fig.show()
```

D:\programs\anaconda\lib\site-packages\sklearn\manifold_t_sne.py:780: FutureWarning:

The default initialization in TSNE will change from 'random' to 'pca' in 1.2.

D:\programs\anaconda\lib\site-packages\sklearn\manifold_t_sne.py:790: FutureWarning:

The default learning rate in TSNE will change from 200.0 to 'auto' in 1.2.

[t-SNE] Computing 151 nearest neighbors...

[t-SNE] Indexed 21094 samples in 0.033s...

[t-SNE] Computed neighbors for 21094 samples in 0.922s...

```

[t-SNE] Computed conditional probabilities for sample 1000 / 21094
[t-SNE] Computed conditional probabilities for sample 2000 / 21094
[t-SNE] Computed conditional probabilities for sample 3000 / 21094
[t-SNE] Computed conditional probabilities for sample 4000 / 21094
[t-SNE] Computed conditional probabilities for sample 5000 / 21094
[t-SNE] Computed conditional probabilities for sample 6000 / 21094
[t-SNE] Computed conditional probabilities for sample 7000 / 21094
[t-SNE] Computed conditional probabilities for sample 8000 / 21094
[t-SNE] Computed conditional probabilities for sample 9000 / 21094
[t-SNE] Computed conditional probabilities for sample 10000 / 21094
[t-SNE] Computed conditional probabilities for sample 11000 / 21094
[t-SNE] Computed conditional probabilities for sample 12000 / 21094
[t-SNE] Computed conditional probabilities for sample 13000 / 21094
[t-SNE] Computed conditional probabilities for sample 14000 / 21094
[t-SNE] Computed conditional probabilities for sample 15000 / 21094
[t-SNE] Computed conditional probabilities for sample 16000 / 21094
[t-SNE] Computed conditional probabilities for sample 17000 / 21094
[t-SNE] Computed conditional probabilities for sample 18000 / 21094
[t-SNE] Computed conditional probabilities for sample 19000 / 21094
[t-SNE] Computed conditional probabilities for sample 20000 / 21094
[t-SNE] Computed conditional probabilities for sample 21000 / 21094
[t-SNE] Computed conditional probabilities for sample 21094 / 21094
[t-SNE] Mean sigma: 1.280513
[t-SNE] KL divergence after 250 iterations with early exaggeration: 70.387543
[t-SNE] KL divergence after 1000 iterations: 0.805038

```

Cluster Of Activities

color

- cycling
- walking
- standing
- downstairs
- upstairs
- sitting

3d Scatterplot in xyz

The scatterplot can properly define which subject did which activity with the use of clusters.

4. Prepare the data to better expose the underlying data patterns to Machine Learning algorithms

prepare your data, is it normalized? are there outlier? Make a training and a test set.

```
In [9]: test_data=pd.read_csv("test.csv", usecols=range(0,7))
train_data=pd.read_csv("train.csv", usecols=range(0,7))
data_data = pd.concat([test_data, train_data])

X = data_data
X = X.dropna()
print(X)
Y = data.activity
Y = Y.dropna()
print(Y)
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=No

print("There are sampels and dimensions for the features", X.shape)
print("There are sampels and dimensions for the Targets", Y.shape)

print("Training Data input")
print(X_train)
print("")
print("Training Activity Data")
print(y_train)
print("")
print("Testing Data input")
print(X_test)
print("")
print("Testing Activity Data")
print(y_test)
```

	Time (s)	Acceleration x (m/s^2)	Acceleration y (m/s^2)	\
0	0.152	-2.59	4.11	
1	0.202	-1.87	4.41	
2	0.252	-2.03	3.78	
3	0.302	-1.83	3.82	
4	0.352	-1.17	2.31	
...	
21089	105.000	8.08	-1.46	
21090	105.000	8.09	-1.46	
21091	105.000	8.09	-1.47	
21092	105.000	8.08	-1.46	
21093	105.000	8.07	-1.45	

	Acceleration z (m/s^2)	Gyroscope x (rad/s)	Gyroscope y (rad/s)	\
0	8.79	-0.525000	-0.393000	
1	8.58	-0.645000	-0.235000	
2	8.92	-0.765000	-0.076200	
3	9.06	-0.320000	-0.044600	
4	8.81	-0.328000	-0.148000	
...	
21089	5.14	-0.000200	0.001260	
21090	5.12	-0.002460	-0.000333	
21091	5.13	-0.001860	0.001330	
21092	5.14	0.000799	0.000133	
21093	5.15	0.002930	0.000732	

Gyroscope z (rad/s)

```

0          0.08300
1          0.29200
2          0.38700
3          0.68200
4          0.50100
...
21089      0.00153
21090     -0.00233
21091     -0.00413
21092     -0.00346
21093     -0.00226

```

[42379 rows x 7 columns]

```

0      cycling
1      cycling
2      cycling
3      cycling
4      cycling

```

```

...
21089   sitting
21090   sitting
21091   sitting
21092   sitting
21093   sitting

```

Name: activity, Length: 42379, dtype: object

There are sampels and dimensions for the features (42379, 7)

There are sampels and dimensions for the Targets (42379,)

Training Data input

	Time (s)	Acceleration x (m/s^2)	Acceleration y (m/s^2)	\
19748	37.1	8.06	-1.480	
3979	199.0	4.28	5.730	
12792	26.6	-3.96	-10.700	
20451	72.8	8.09	-1.480	
18621	227.0	-3.81	-0.734	
...	
20266	63.6	8.08	-1.490	
1708	85.6	5.08	1.220	
19440	21.4	8.08	-1.470	
19189	255.0	-3.09	-1.010	
10847	37.4	10.40	5.440	

	Acceleration z (m/s^2)	Gyroscope x (rad/s)	Gyroscope y (rad/s)	\
19748	5.150	-0.000200	0.00160	
3979	-9.730	-3.110000	-0.53200	
12792	0.426	1.040000	1.18000	
20451	5.130	-0.001260	0.00186	
18621	-8.950	-0.000342	0.02890	
...	
20266	5.120	0.000333	0.00020	
1708	-7.410	-2.430000	-0.69700	
19440	5.160	0.002730	-0.00280	
19189	-9.200	-0.032400	-0.00764	
10847	3.790	0.424000	0.22200	

	Gyroscope z (rad/s)
19748	-0.000599
3979	-1.670000
12792	0.733000
20451	0.000399
18621	-0.028400
...	...
20266	-0.000799
1708	-0.991000
19440	0.000533
19189	-0.143000
10847	-0.410000

[33903 rows x 7 columns]

Training Activity Data

19748 sitting
3979 cycling
12792 walking
20451 sitting
18621 sitting

...

20266 sitting
1708 cycling
19440 sitting
19189 sitting
10847 cycling

Name: activity, Length: 33903, dtype: object

Testing Data input

	Time (s)	Acceleration x (m/s ²)	Acceleration y (m/s ²)	\
18519	45.80	4.950	7.23	
6583	4.55	-1.050	-4.06	
10589	79.20	0.745	9.86	
15858	75.30	-0.439	9.64	
12362	168.00	1.940	9.49	
...	
18160	27.80	4.880	7.38	
6141	307.00	1.790	7.25	
1566	78.50	0.674	9.71	
2942	147.00	3.050	6.48	
9813	40.40	2.940	9.49	

	Acceleration z (m/s ²)	Gyroscope x (rad/s)	Gyroscope y (rad/s)	\
18519	4.1600	-0.13700	-4.93000	
6583	-2.6500	-0.76600	-0.68900	
10589	-0.0198	-0.01960	0.07310	
15858	0.2140	-0.00126	-0.00626	
12362	-1.5000	0.01500	0.10600	
...	
18160	3.4600	-0.00919	0.04000	
6141	-6.8100	1.38000	1.05000	
1566	-4.9100	-0.61500	-0.45800	
2942	-7.7600	2.52000	0.63400	
9813	-0.0533	-0.05860	-0.36800	

	Gyroscope z (rad/s)
18519	-4.24000
6583	-0.03410
10589	-0.00777
15858	-0.00739
12362	-0.04330
...	...
18160	-0.03000
6141	0.50000
1566	-0.06290
2942	1.25000
9813	0.05130

[8476 rows x 7 columns]

Testing Activity Data

18519 upstairs
6583 upstairs
10589 standing
15858 standing
12362 standing

...

```
18160    upstairs
6141     cycling
1566     cycling
2942     cycling
9813     standing
Name: activity, Length: 8476, dtype: object
```

Data is getting shuffled meaning the rows are getting randomized. Splitting data in data inputs and data outputs (activity) which is used for learning. 2 of the 3 subjects are used to train the algorithms and the rest is used to test the algorithm.

5. Explore many different models and short-list the best ones.

Explore / train and list the top 3 algorithms that score best on this dataset.

```
In [10]: #algorithm functions
print("start algo")
#-----
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)
print("knn")
#-----
svclassifier_pol = SVC(kernel='poly', degree=8)
svclassifier_pol.fit(X_train, y_train)
y_pred_pol = svclassifier_pol.predict(X_test)
print("pol")
#-----
svclassifier_rbf = SVC(kernel='rbf')
svclassifier_rbf.fit(X_train, y_train)
y_pred_rbf = svclassifier_rbf.predict(X_test)
print("rbf")
#-----
dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
y_pred_dtc = dtc.predict(X_test)
print("dtc")
#-----
etc = ExtraTreesClassifier()
etc.fit(X_train, y_train)
y_pred_etc = etc.predict(X_test)
print("etc")
#-----
rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
y_pred_rfc = rfc.predict(X_test)
print("rfc")
#-----
gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred_gnb = gnb.predict(X_test)
print("gnb")
#-----
qda = QuadraticDiscriminantAnalysis()
qda.fit(X_train, y_train)
y_pred_qda = qda.predict(X_test)
print("qda")
```

```

#plotting non normalized confusion matrices
print("KNN Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_knn))
print("Poly Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_pol))
print("Rbf Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_rbf))
print("DecisionTree Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_dtc))
print("ExtraTrees Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_etc))
print("RandomForest Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_rfc))
print("GaussianNB Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_gnb))
print("QuadraticDiscriminantAnalysis Not normalized confusion matrix")
print(confusion_matrix(y_test,y_pred_qda))


#lists of algorithms
class_names = data.activity
titles_options = [
    ("KNN Normalized confusion matrix", "true"),
    ("Poly_SVM Normalized confusion matrix", "true"),
    ("rbf_SVM Normalized confusion matrix", "true"),
    ("DecisionTree Normalized confusion matrix", "true"),
    ("ExtraTrees Normalized confusion matrix", "true"),
    ("RandomForest Normalized confusion matrix", "true"),
    ("GaussianNB Normalized confusion matrix", "true"),
    ("QuadraticDiscriminantAnalysis Normalized confusion matrix", "true")
]

algo_names = [
    knn,
    svcclassifier_pol,
    svcclassifier_rbf,
    dtc,
    etc,
    rfc,
    gnb,
    qda
]

title_names = [
    "knn",
    "pol",
    "rbf",
    "dtc",
    "etc",
    "rfc",
    "gnb",
    "qda"
]

pred_names = [
    y_pred_knn,
    y_pred_pol,
    y_pred_rbf,
    y_pred_dtc,
    y_pred_etc,
    y_pred_rfc,
    y_pred_gnb,
    y_pred_qda
]

```

```

]

#plotting normalized matrices
i = 0
for title, normalize in titles_options:
    disp = ConfusionMatrixDisplay.from_estimator(algo_names[i], X_test, y_test, cmap=plt
    disp.ax_.set_title(title)
    i = i + 1
plt.show()

#printing and calculating accuracy
i = 0
for title in title_names:
    print(title, "accuracy :", round(sklearn.metrics.accuracy_score(y_test, pred_names[i]
    i = i + 1

```

start algo

D:\programs\anaconda\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

knn

pol

rbf

dtc

etc

rfc

gnb

qda

KNN Not normalized confusion matrix

```

[[2250   7   7  17   2   9]
 [  33 396   4  36  59  80]
 [   1   0 2218   0   0   0]
 [   5   2   0 1398   3   9]
 [  12  72   8  34 569  72]
 [  44  66   2  74  38 949]]

```

Poly Not normalized confusion matrix

```

[[1243   0  997  52   0   0]
 [   0   0  608   0   0   0]
 [  10   0 2209   0   0   0]
 [   0   0 1197 220   0   0]
 [   0   0  767   0   0   0]
 [   0   0 1173   0   0   0]]

```

Rbf Not normalized confusion matrix

```

[[2020   2  167 102   0   1]
 [  130  65  49 191   6 167]
 [   8   3 1875  13   0 320]
 [   3   1  10 1158   0 245]
 [  73  56  70 175 165 228]
 [ 114   2 116 352   0 589]]

```

DecisionTree Not normalized confusion matrix

```

[[2203  28   7   7  12  35]
 [  36 347   1   9 121  94]
 [   2   2 2212   0   1   2]
 [  14   8   0 1370  12  13]
 [  13 115   3  10 551  75]
 [  27  72   0  15  75 984]]

```

ExtraTrees Not normalized confusion matrix

```

[[2256  15   0   5   5  11]
 [  23 448   0   6  61  70]]

```



```

[ 1 2 2216 0 0 0]
[ 2 1 0 1392 5 17]
[ 10 59 1 4 633 60]
[ 7 21 0 2 14 1129]]
RandomForest Not normalized confusion matrix
[[2242 24 0 6 7 13]
 [ 14 429 2 4 85 74]
 [ 2 2 2215 0 0 0]
 [ 2 6 0 1387 4 18]
 [ 8 62 0 2 644 51]
 [ 7 29 0 1 19 1117]]

```

```

GaussianNB Not normalized confusion matrix
[[1962 195 49 31 9 46]
 [ 55 231 4 56 66 196]
 [ 40 3 2165 4 1 6]
 [ 0 5 0 1385 5 22]
 [ 22 139 3 37 285 281]
 [ 70 126 8 118 48 803]]

```



```



QuadraticDiscriminantAnalysis Not normalized confusion matrix
[[1980 115 39 25 18 115]
 [ 56 264 2 34 56 196]
 [ 34 6 2160 4 1 14]
 [ 1 12 0 1381 4 19]
 [ 29 165 2 18 288 265]
 [ 45 72 7 110 21 918]]

```

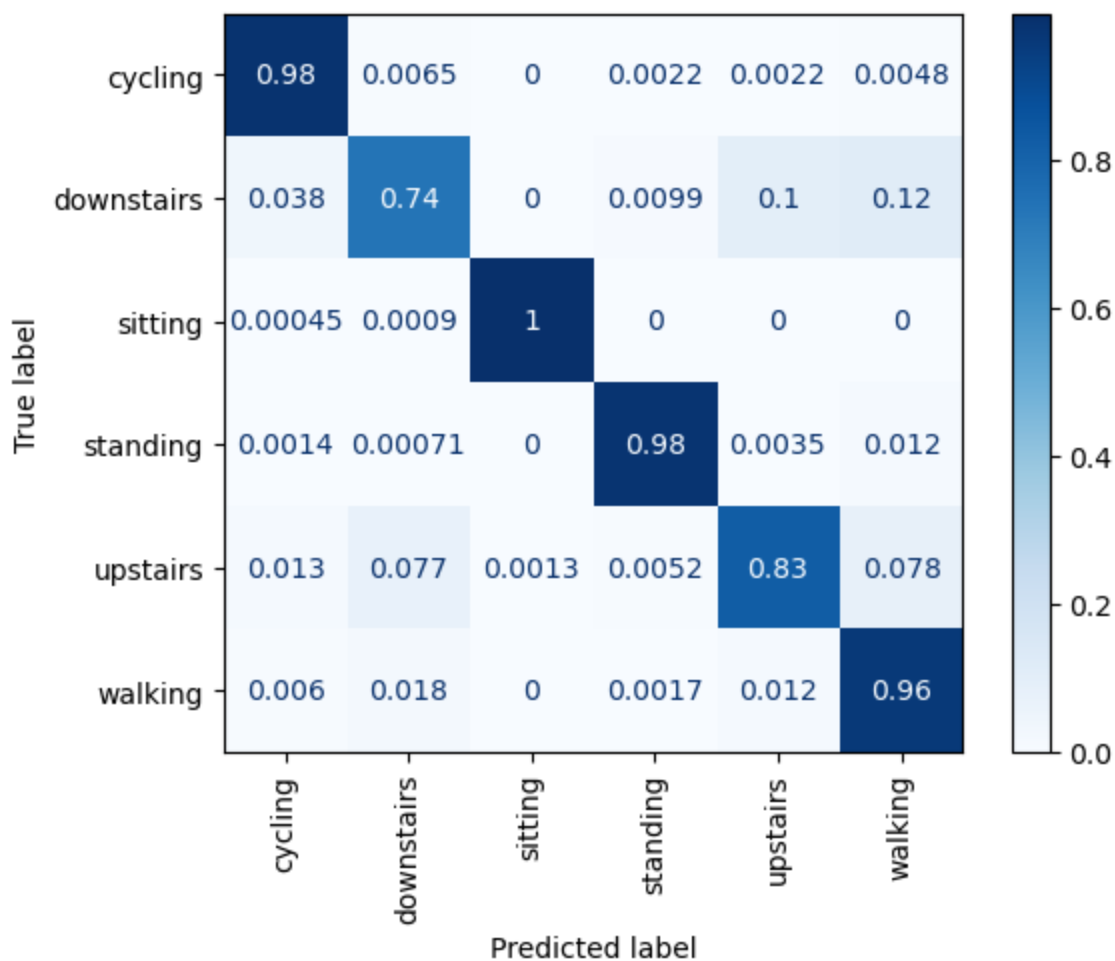
D:\programs\anaconda\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. ``skew``, ``kurtosis``), the default behavior of ``mode`` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of ``keepdims`` will become `False`, the ``axis`` over which the statistic is taken will be eliminated, and the value `None` will no longer be accepted. Set ``keepdims`` to `True` or `False` to avoid this warning.








ExtraTrees Normalized confusion matrix



RandomForest Normalized confusion matrix





knn accuracy : 91.79 %

```

pol accuracy : 43.32 %
rbf accuracy : 69.28 %
dtc accuracy : 90.46 %
etc accuracy : 95.26 %
rfc accuracy : 94.79 %
gnb accuracy : 80.59 %
qda accuracy : 82.48 %

```

Knn with k is 5. The subject with the highest correlation are eliminated. the perfect model should have a score of 1 when all predictions are correct.

here we can see what algorithms perform the best with our data.

The best three we have chosen to build into our model.

These are the:

Knn algorithm

Etc algorithm

Rfc algorithm

6. Fine-tune your models and combine them into a great solution.

can you get better performance within a model? e.g if you use a KNN classifier how does it behave if you change K (k=3 vs k=5 vs k=?). Which parameters are here to tune in the chosen models?

```

In [11]: X_train, X_test, y_train, y_test = train_test_split( X,Y, stratify=Y, test_size=0.2, ran

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_lin = knn.predict(X_test)

y_train_pred = knn.predict(X_train)
y_test_pred = knn.predict(X_test)

knn_train_accuracy = accuracy_score(y_train, y_train_pred)

# Test set performance
knn_test_accuracy = accuracy_score(y_test,y_test_pred)

print('Model performance for Training set')
print('- Accuracy: %s' % round(knn_train_accuracy*100,2), "%")

print('-----')
print('Model performance for Test set')
print('- Accuracy: %s' % round(knn_test_accuracy*100,2), "%")

```

D:\programs\anaconda\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

D:\programs\anaconda\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
Model performance for Training set
- Accuracy: 94.83 %
-----
```

```
Model performance for Test set
- Accuracy: 91.84 %
```

D:\programs\anaconda\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning:

Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

The knn algorithm was trained and was tested with an accuracy of 91.84%

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(
        X, Y, stratify=Y, test_size=0.2, random_state=42
    )

    rfc = RandomForestClassifier()
    rfc.fit(X_train, y_train)
    y_pred_pol = rfc.predict(X_test)

    y_train_pred = rfc.predict(X_train)
    y_test_pred = rfc.predict(X_test)

    rfc_train_accuracy = accuracy_score(y_train, y_train_pred)

    # Test set performance
    rfc_test_accuracy = accuracy_score(y_test, y_test_pred)

    print('Model performance for Training set')
    print('- Accuracy: %s' % round(rfc_train_accuracy*100,2), "%")

    print('-----')
    print('Model performance for Test set')
    print('- Accuracy: %s' % round(rfc_test_accuracy*100,2), "%")
```

```
Model performance for Training set
- Accuracy: 100.0 %
-----
```

```
Model performance for Test set
- Accuracy: 94.84 %
```

The Rfc algorithm was trained and was tested with an accuracy of 94.84%

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(
        X, Y, stratify=Y, test_size=0.2, random_state=42
    )

    etc = ExtraTreesClassifier(n_estimators=500, min_samples_split=3, max_depth=500)
```

```

etc.fit(X_train, y_train)
y_pred_etc = etc.predict(X_test)

y_train_pred = etc.predict(X_train)
y_test_pred = etc.predict(X_test)

etc_train_accuracy = accuracy_score(y_train, y_train_pred)

# Test set performance
etc_test_accuracy = accuracy_score(y_test, y_test_pred)

print('Model performance for Training set')
print('- Accuracy: %s' % round(etc_train_accuracy*100,2), "%")

print('-----')
print('Model performance for Test set')
print('- Accuracy: %s' % round(etc_test_accuracy*100,2), "%")

```

```

Model performance for Training set
- Accuracy: 100.0 %
-----
Model performance for Test set
- Accuracy: 95.38 %

```

The Etc algorithm was trained and was tested with an accuracy of 95.38%

```

In [14]: estimator_list = [
          ('knn', knn),
          ('rfc', rfc),
          ('etc', etc) ]
stack_model = StackingClassifier(
    estimators=estimator_list, final_estimator=LogisticRegression()
)
stack_model.fit(X_train, y_train)
y_train_pred = stack_model.predict(X_train)
y_test_pred = stack_model.predict(X_test)

stack_model_train_accuracy = accuracy_score(y_train, y_train_pred)
stack_model_test_accuracy = accuracy_score(y_test, y_test_pred)

print('Model performance for Training set')
print('- Accuracy: %s' % round(stack_model_train_accuracy*100,2), "%")

print('-----')
print('Model performance for Test set')
print('- Accuracy: %s' % round(stack_model_test_accuracy*100,2), "%")

```

D:\programs\anaconda\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning:

```

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

Model performance for Training set
- Accuracy: 100.0 %

```

Model performance for Test set
- Accuracy: 95.85 %

Here we build a model from our three algorithms. We combined the and made one model. This model has an accuracy of 95,75%. This is better than the best single algorithm and proves that the model is a good solution for the algorithms.

7. Present your solution.

Explain why you would choose for a specific model

Total code of the assignment is shown above

8. Launch, monitor, and maintain your system.

Can you Deployment the model?

NOTE: The app provides the option for remote access, so you are able to get live sensordata from the phone

9. Additional Questions

1. Explain the chosen motions you chose to be classified.
 2. Which of these motions is easier/harder to classify and why?
 3. After your experience, which extra sensor data might help getting a better classifier and why?
 4. Explain why you think that your chosen algorithm outperforms the rest?
 5. While recording the same motions with the same sensor data, what do you think will help improving the performance of your models?
-
1. these particular motions are based on our daily activities wich means it is easy to gather a lot of data
 2. by looking at the graph in part 3 we can conclude that both walking upstairs and downstairs are very hard to tell apart.
 3. some way of measuring the altitude accurately can be used to dertermine if the person is walking upstairs or downstairs.
 4. in part 5 there are 8 diffrent algorithms tested. by looking at the results we determined the 3 best ones. these were later combined into 1 model with a high accuracy.
 5. the more data is accuired, the more accurate the model will become. it will also help if the same motions are recorded on diffrent people to have more variety in the data.